



**WPI**

# **Final Project – Monocular Depth Estimation**

Pranay Katyal, Anirudh Ramanathan

# Problem Statement

Monocular Depth Estimation from Drone Camera Images



# Overview

---

- The goal of this project is to estimate depth from images captured by a monocular camera onboard a drone. Monocular depth estimation is a challenging problem in computer vision, which this project will address using deep learning methods.
- Reproduce the results from the paper:
  - M. Folder et al., “Parallax Inference for Robust Temporal Monocular Depth Estimation in Unstructured Environments”, 2022.
    - This technique is unique because it **fuses sequential camera images** with corresponding **camera translation and rotation measurements (from GPS/IMU)** to produce **metric depth estimates**. The authors have open-sourced their code, which is implemented in TensorFlow.

# Goals

---

- The project has two main phases:
  - **Synthetic Data Training:** We will first follow the instructions in the official GitHub repository to train the network on the **MidAir dataset**, a **synthetic dataset**. The objective is to **replicate the performance** reported in the paper.
    - <https://midair.ulg.ac.be/>
  - **Real-World Data Fine-Tuning:** Next, we will use the **pre-trained model** and **fine-tune** it on the **UseGeo dataset**, which consists of **real-world images**. The final objective is to predict depth for images in the UseGeo dataset and compare your predictions to the ground-truth depth maps.
    - <https://github.com/3DOM-FBK/UseGeo/tree/master>

# Dataset - MidAir Synthetic Dataset

---

- **Mid-Air**, The *Montefiore Institute Dataset of Aerial Images and Records*, is a multi-purpose synthetic dataset for low altitude drone flights. It provides a large amount of synchronized data corresponding to flight records for multi-modal vision sensors and navigation sensors mounted on board of a flying quadcopter. They also provide Train and Test Data.
- Due to the sheer size of their dataset, we had to select a fragment (**Trajectory0000**) one we want to download. Their download procedure works as follow:
  - Select the desired visual sensors;
  - Select the desired training data, *i.e.* trajectories and climate setups;
  - Select desired benchmarks;
  - Submit the form.
- Image Type Options:
  - Left RGB, Right RGB, **Down RGB**, Surface Normals, **Depth**, **Stereo Disparity**, Stereo Occlusion, Semantic Segmentations.



# Dataset - UseGeo Real Dataset

- **UseGeo**, is the new large-scale real-world set of data Depth estimation, Multiview 3D Reconstruction, etc.
- Number of Available Datasets
  - **3 Total Datasets** available.
  - Approximately **100GB** worth of Data.
  - **Images, Depths**, Dense point clouds.
  - **Image Size** : 7952 x 5304 px
  - **Camera Poses** : X0,Y0, Z0, Omega, Phi, Kappa, **Camera Intrinsics**.
- Flight Specifications
  - **Area Covered**: 1100 × 650 meters
  - **Average Flight Height**: 80 m above ground
  - **GSD (Ground Sampling Distance)**: approx. 2 cm
  - **Image Overlap**: 80% (forward) – 60% (side)
  - **LiDAR Point Cloud Density**: ~50 points/m<sup>2</sup>



# Dataset - Summary

---


- **MidAir:** 550 synthetic aerial images, relative poses
  - Gained from Trajectory0000 Alone!
- **UseGeo:** 828 real drone images, absolute GPS poses
  - All 3 Datasets on GitHub combined.
- **Train/Validation Split:** 80/20 for UseGeo
  - We split across each Dataset rather than a 2:1 Training and Eval split.

# M4Depth Architecture

---

- **Feature Pyramid Encoder**
  - Extracts hierarchical features at 6 resolution levels
  - Uses Domain-Invariant Normalization (DINL)
- **Temporal Depth Estimator**
  - **Temporal Recurrence**: Warps features from previous frames using camera motion.
  - **Cost Volume**: Correlates features across frames to find correspondences.
  - **Depth Refinement**: Multi-scale processing for accurate predictions
- **Parallax to Depth Conversion**
  - Converts network output (parallax/disparity) to metric depth.
  - **Critical**: Uses camera motion (translation & rotation) via `parallax2depth()`.
  - Formula:  $\text{depth} = f(\text{disparity}, \text{camera\_motion}, \text{intrinsics})$
- **Key Insight**: Model fundamentally depends on accurate pose information!

# M4Depth Architecture



# Phase 1 – Results for Pretrained Model Performance on MidAir


---

- Evaluation Setup:
  - Pretrained M4Depth model (71 epochs on MidAir training data)
  - Evaluated on MidAir trajectory0000 validation set (550 samples)
  - Same domain as training (synthetic aerial imagery)
- Performance Metrics:


Metric	Value	Interpretation
RMSE	2.65m	Average depth error
Abs Rel	0.021	2.1% relative error
$\delta < 1.25$	98.6%	Near-perfect accuracy

# Phase 1 – Results for Pretrained Model Performance on MidAir




- Training Plot:






- Phase 1 Evaluation:



# Phase 1 – Results - 2 of the 10 Samples



- Each row: RGB | Predicted Depth | Ground Truth
- The base model – from their **Checkpoint 71**
- Good performance!



# Phase 2 – Results for Training on Real-World UseGeo Data

---

- Evaluation Setup:
  - Custom DataLoader for UseGeo (828 real drone images)
  - 80/20 split: 661 train, 167 validation
  - 146 epochs, ~3 hours training
  - **Note:** Pretrained checkpoint failed to load → trained from scratch

- Observations:
  - Dramatic metric improvement
  - Training converged smoothly (see curves)
  - **BUT:** Visual inspection revealed problems...

- Performance Metrics:


Metric	Pretrained (MidAir)	Fine-tuned (UseGeo)	Improvement
RMSE (m)	2.6543	0.6201	<b>76.6% ↓</b>
Abs Rel	0.0212	0.0031	<b>85.4% ↓</b>
$\delta < 1.25$ (%)	98.59	99.80	<b>1.21% ↑</b>
$\delta < 1.25^2$ (%)	99.28	99.95	<b>0.67% ↑</b>
$\delta < 1.25^3$ (%)	99.54	100.00	<b>0.46% ↑</b>

# Phase 2 – Results for Training on Real-World UseGeo Data


- Training Plot:



- Phase 2 Evaluation:



# Phase 2 – Results - 2 of the 10 Samples



- Each row: RGB | Predicted Depth | Ground Truth
- Tuning didn't work so we trained from scratch using 80-20 split.
- Model **Checkpoint 145** used
- Okayish performance!
  - 2 main reasons
    - Low amount of Data
    - Model Architecture issues possibly.

# The Problem: Pose Format Incompatibility

---

- What We Discovered:
  - Depth predictions were **200 million meters** instead of 60-140m!
- Root Cause: Pose Representation Mismatch
- Scale difference:  $\sim 2,000,000 \times$
- Impact on M4Depth:
  - `parallax2depth()` expects relative motion (meters)
  - Received absolute GPS coordinates (hundreds of kilometers!)
  - Internal calculations:  $\text{depth} = f(\text{disparity}, \text{translation})$
- Result: Predictions in millions of meters

Dataset	Pose Format	Translation Values
MidAir	Relative motion	0-20 meters
UseGeo	Absolute GPS	$\sim 498,340$ meters

# The Paradox: Good Metrics, Bad Predictions

---

- Metrics Computed on Wrong Tensors
  - Evaluated on **internal network representations**
  - NOT on final metric depth after `parallax2depth()`
  - Network outputs had correct *relative patterns*.
- Training in Log-Space
  - Loss:  $\log(200M) - \log(100) = \text{finite, trainable gradient}$
  - Model learned consistent patterns despite wrong scale
  - Log-space masks absolute magnitude errors
- Pattern Matching Without Scale
  - Network learned: "building closer than Roads"
  - But absolute values: millions of meters!
  - Relative ordering correct, absolute scale wrong

# Debugging Attempts: What We Tried

---

- **Solution 1: Pose Conversion Tool**

- Created `convert_usegeo_poses.py`
  - Converted absolute GPS → relative motion
- Algorithm:  $t_{\text{rel}} = R_{\text{prev}}^T \times (t_{\text{curr}} - t_{\text{prev}})$
- **Result:** Mean 16.7m, median 12.8m (correct range!)

- **Solution 2: Fine-tune with Relative Poses**

- Loaded MidAir checkpoint + relative pose CSVs
- **Result:** Training crashed at batch 47 with NaN loss

- **Solution 3: Lower LR + Gradient Clipping**

- Learning rate: 0.0001 → 0.00001 ( $10 \times$  lower), and Added gradient clipping (`clipnorm=1.0`)
- **Result:** Still crashed at batch 47 with NaN

- **Solution 4: Train from Scratch with Relative Poses**

- Violates assignment: must use pretrained weights
- **Result:** Doesn't work well due to lack of data 350GB.

# Lessons learnt

---

- Transfer Learning Requires Complete Data Compatibility
  - Not just image domain adaptation
  - ALL input representations must match
  - Pose format as critical as visual features.
  - Hidden dependencies in pretrained weights
- Visual Validation is Essential
  - Metrics can mask fundamental failures
  - Always inspect actual outputs, not just numbers
  - Check value ranges for physical plausibility
- Systematic Debugging Reveals Root Causes
  - Traced through: metrics → depth values → pose format
  - Multiple solution attempts confirmed hypothesis

# Conclusions & Future Directions

---

- Does pre-training on synthetic MidAir help, hurt, or have no effect on real UseGeo images?
  - We Cannot Definitively Conclude Transfer Learning Effectiveness
    - Pose format incompatibility **prevented** proper transfer
    - Phase 2 trained **from scratch**, not fine-tuned
  - What We Can Say:
    - **Pose representation compatibility** is a prerequisite for transfer
    - **Transfer learning requires** matching ALL input formats, not just images
    - **From-scratch training worked**, showing architecture is sound
- **Expectation:**
  - IF pose formats matched → pretrained weights likely would help (Similar scene types, motion patterns, depth estimation task) just like it did when we did Transfer Learning on CNN images.



# WPI

**Thank You!**

