

Machine Learning for Robotics: **Introduction**

Prof. Navid Dadkhah Tehrani



Artificial Intelligence (AI):

Born at a workshop at Dartmouth college in 1956. It is simulation of human intelligence processes by machines, especially computer systems, enabling them to perform tasks such as learning, reasoning, problem-solving, and understanding natural language.

Machine Learning (ML):


deals with the statistical inference problem of finding a predictive function based on data.

Deep learning :

make use of the deep neural network as opposed to using classical machine learning algorithms (e.g. k-mean clustering, decision tree, linear regression, ...)

Artificial General Intelligence (AGI):

Multi-purpose AI mimicking human intelligence across a variety of tasks as opposed to narrow AI that's only specialized in one task.



Why Deep Learning?

The performance of the deep learning models get better as we feed them more data.




Image from: C. Aggarwal, Neural Networks and Deep learning

Some applications of Deep Learning in Robotics

- Image classification
- Image semantic segmentation
- Depth estimation from camera
- Object detection.
- Object tracking.
- Model prediction
- Reinforcement learning
- Imitation Learning
- Fault/Anomaly detection :
can be employed for detecting anomalies or faults in robotic systems
by capturing the normal relationships and dependencies within the system
and identifying deviations from the norm.

Non robotics applications:

- Email spam detection
- Fingerprint identification
- Stock market prediction
- Face detection
- Reading checks at bank's ATM machine
- Amazon's product recommendation
- Netflix movie recommendation
- medical diagnosis
- Siri or Alexa speech recognition
- Large Language models
- Text translation

Datasets

Researchers in each field in deep learning make their own dataset and often time release it to the public so more researcher can develop algorithms on the data set.

Some of the famous datasets:

(pretty much a dataset exist nowadays for anything you want to do!).

- ImageNet: 14 million images, annotated in 20,000 categories
<https://www.image-net.org/>
- CIFAR-10 (CIFAR-100): consists of 60k, 32x32 color images in 10 (100) classes.
<https://www.cs.toronto.edu/~kriz/cifar.html>
- COCO: 330K images, 1.5M objects, annotated in 91 categories
<https://cocodataset.org/>
- nuScenes: 1,000 driving scenes, containing 1.4M images, 1.4M bounding boxes and 390K lidar sweeps
<https://www.nuscenes.org/nuscenes>
- Kitti: 100,000 images across multiple hours of driving. 7.5K in object detection benchmark
<https://www.cvlibs.net/datasets/kitti/>

Synthetic datasets:

They are generated by running high fidelity simulation of vehicles with high quality graphic engines such as Unreal engine or Unity.

Mid-Air dataset: 79 minutes of drone flight in simulator with stereo camera and IMU/GPS.

<https://midair.ulg.ac.be/>

Reading Material for the Course:


This course focuses 100 percent on deep learning. We do not cover classical machine learning.

Textbook: Deep dive into deep learning, by A. Zhang, Z. Lipton, M. Li , A. Smola, 2023

Also available online via <https://d2l.ai/>

The book covers the basic material that are prerequisites for advanced concepts.

For robotics application we mainly use research paper.



Tools used in this course:

- Python
 - Pytorch → <https://pytorch.org/tutorials/>
 - Numpy
 - Matplotlib
-



Released in 2017

Pytorch API is closer to Numpy. Based on trend in the recent research papers, it seems that the deep learning community likes Pytorch better.

Recommended IDE:



Visual Studio Code

multi-platform, multi-language support: Python, Julia, C/C++, etc,
powerful debugger.

Tensor

Tensors are basically high dimensional arrays.

In deep learning everything is done with tensors.

All the data is first converted to tensors, and then goes into a neural networks.

Neural network itself is of collection of tensors with various operation such addition, multiplication, normalization, nonlinearity,

The output is a tensor/tensors of desired output.

Q: why can't we do all of this in numpy?

A: Pytorch has GPU support, which means that we can load the data and neural network in the GPU and take advantage of parallel processing.

GPUs have parallelization for linear algebra operations.

Also Pytorch has automatic differentiations which is used in training of neural networks.

Q: Since Pytorch is in in Python, isn't it too slow?

A: The core functionality of Pytorch is written in C++, with python binding (i.e. C++ backend). And the parts that are in native Python are written very efficiently.

According to a study, Pytorch is only responsible for about 10 percent of the slowdowns. Which means that even if we use a full C++ API we might only get 10 percent performance improvement. This not substantial compared to difficulty of working with C++ API.

What is Cuda:

Cuda is a C++ library for Nvidia GPUs.



Similarly ROCm is GPU library for AMD processors. And arm compute library has GPU support for Arm processors.



Cuda is the most popular one, but other companies are also trying to get into this market.

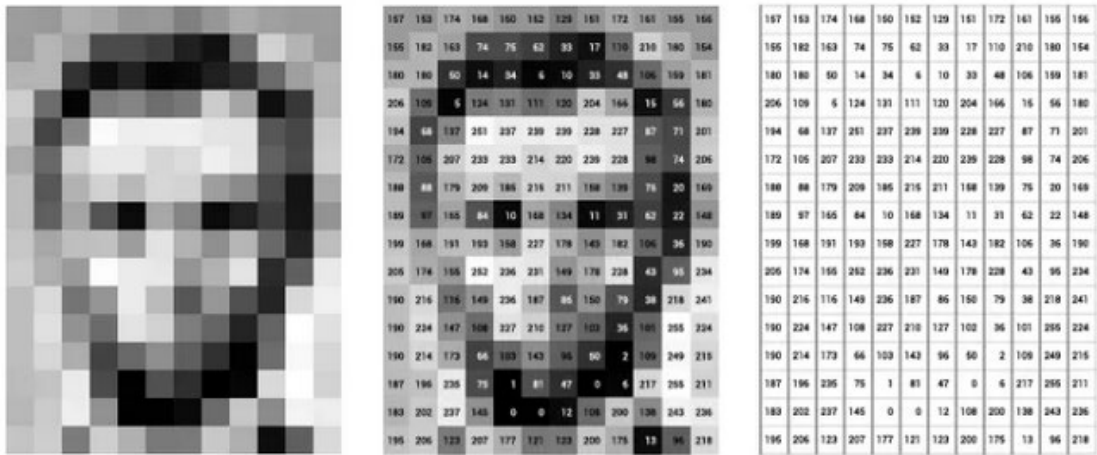


Rank 0 tensor: scalar

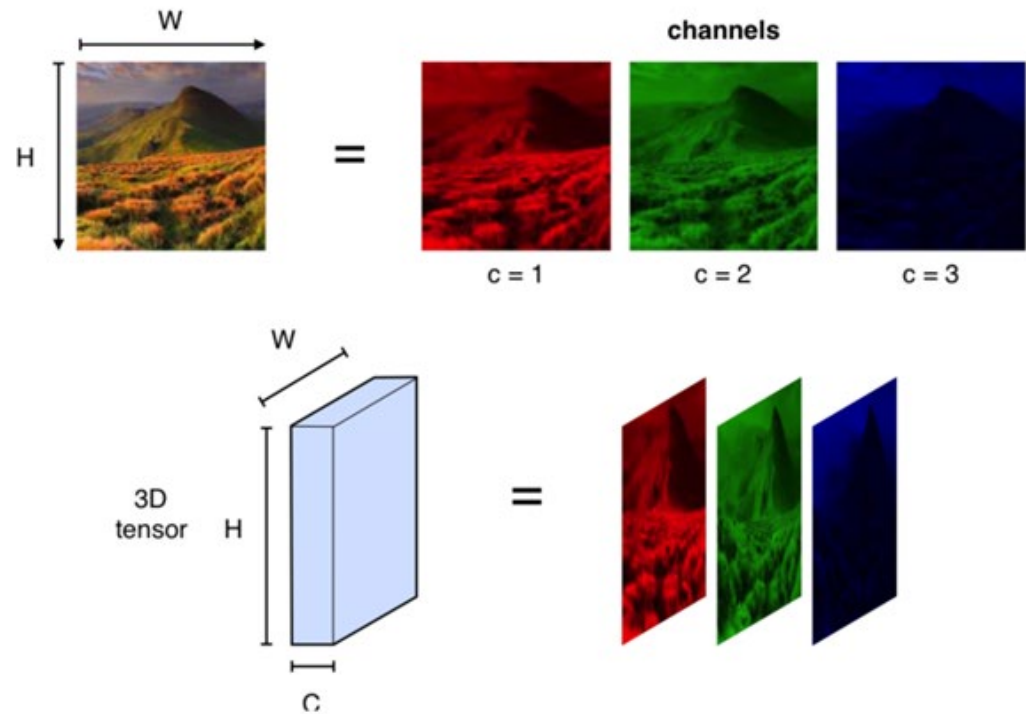
Rank 1 tensor: vector

Rank 2 tensor (2D tensor):

Gray scale image:

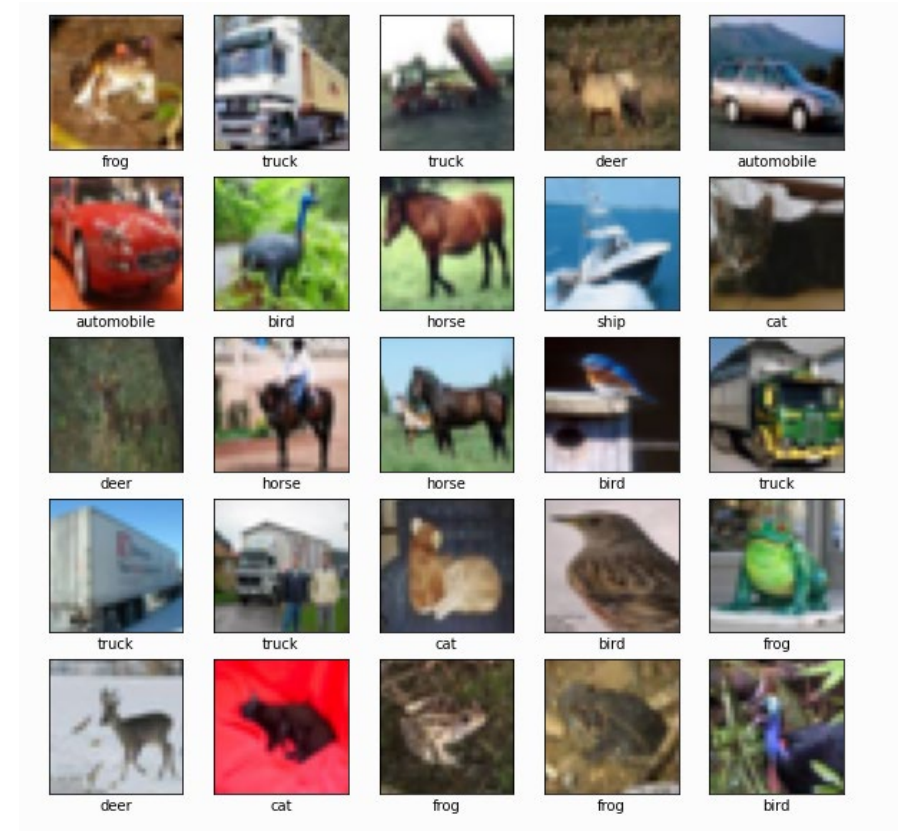


Rank 3 tensors (3D tensors)



Rank 4 Tensor (4D tensor)

Batch of RGB images:



From CIFAR-10 dataset