

# Machine Learning for Robotics: **Classification**

Prof. Navid Dadkhah Tehrani



In regression, we're predicting a float value as output.

In classification, we're predicting a class type that's an integer number. For example, what object is in the image out of 100 possible objects ( $i \in [1, 2, \dots, 100]$ )

We have two problems:

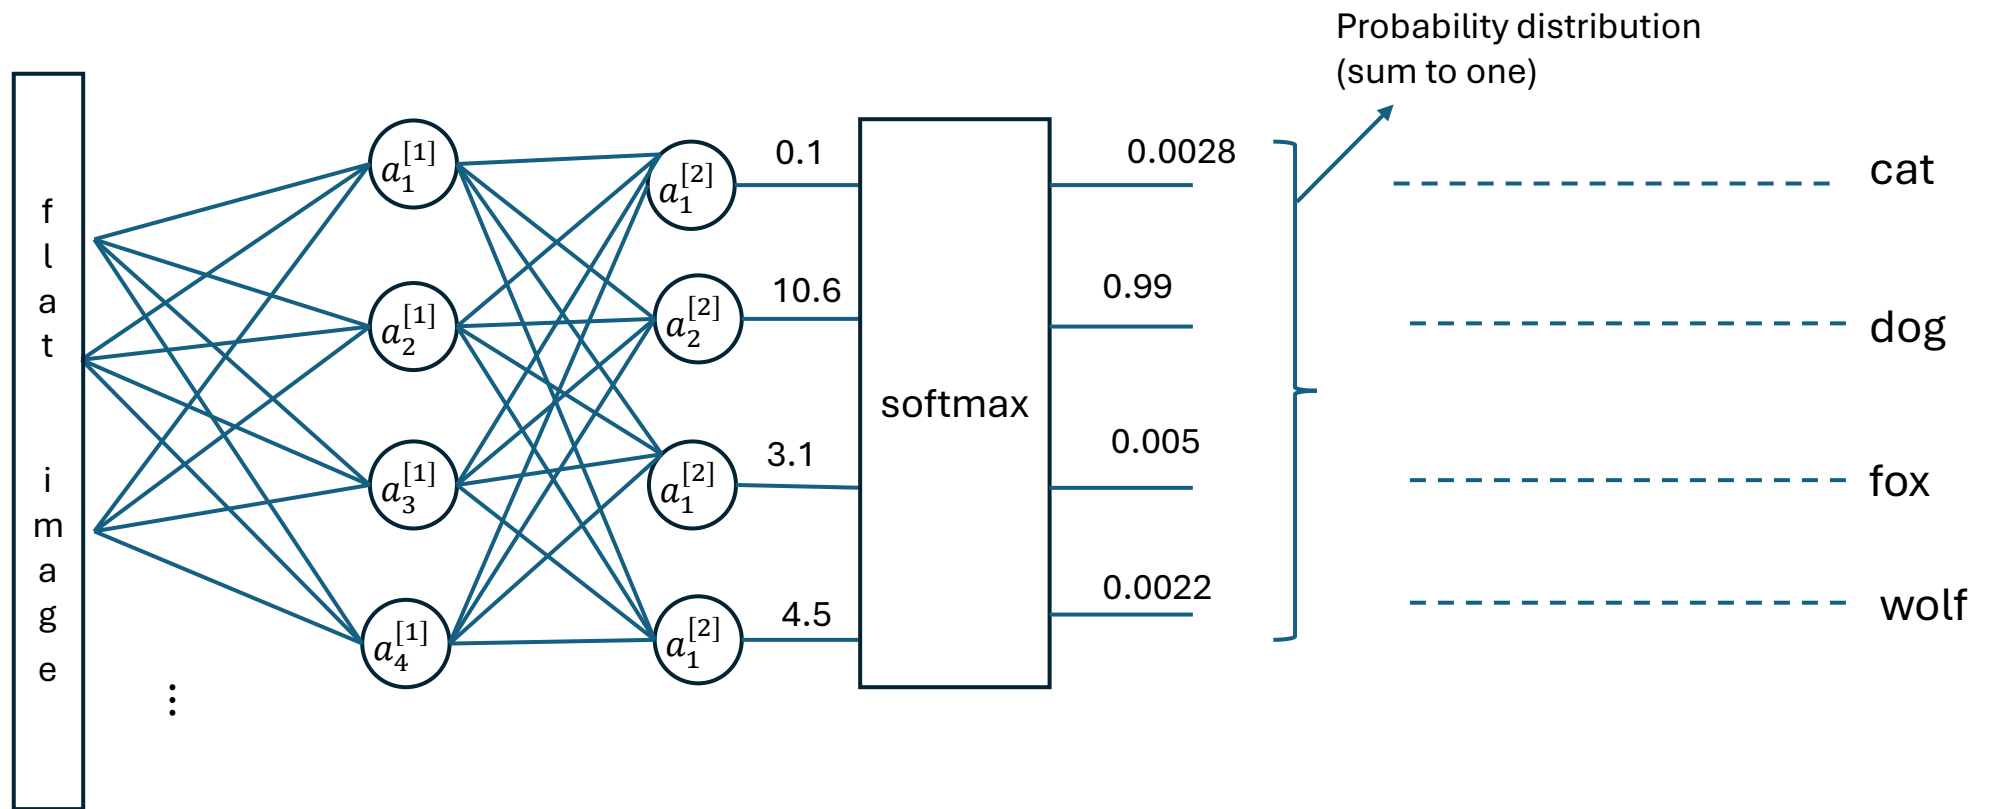
- The output of the neural network is a float number (tensor) not an integer.
- How do we define the continuous and differentiable loss function for classification?

assume that we want to classify an image in 4 categories (cat, dog, fox, wolf)

Say the image is 32-by-32 gray scale. We first flatten the image into a 1024 vector.

Then we need is a neural network that output 4 values.

We then pass the 4 output to a softmax function, to convert the output to a probability distribution.



## Definition of softmax

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{o}) \quad \text{where} \quad \hat{y}_i = \frac{\exp(o_i)}{\sum_j \exp(o_j)}.$$

$$\text{Example: } \mathbf{o} = \begin{bmatrix} 1 \\ 5 \\ 10 \\ 0.1 \end{bmatrix} \rightarrow \hat{\mathbf{y}} = \begin{bmatrix} 0.001 \\ 0.0067 \\ 0.9931 \\ 0.000 \end{bmatrix}$$

Softmax turn a vector into a probability distribution.

Softmax is also a soft max! in other words, it's continuous and differentiable way of taking maximum of bunch of numbers.

One-hot encoding of class labels

before we define the loss function, we must define one-hot encoding of class labels.

	cat class	Dog class	Fox class	Wolf class	Integer class label
Training example #1	1	0	0	0	1
Training example #2	0	0	1	0	3
Training example #3	1	0	0	0	1
Training example #4	0	0	1	0	3
...	0	1	0	0	2
...	1		0	0	1
...	0	0	1	0	3
...	0	0	0	1	4
Training example #n	0	0	0	1	4

## Cross-entropy Loss function

$$L = \sum_{i=1}^n \sum_{j=1}^h -y_j^{[i]} \log(\hat{y}_j^{[i]}) \quad \text{or} \quad L = \sum_{i=1}^n -y_j \cdot \log(\hat{y}_j) \quad ( '.' \text{ is dot product})$$

$n$ : number of training examples

$h$ : number of classes

$\hat{y}_j^{[i]}$ : predicted output of the NN

$y_j^{[i]}$ : onehot encoding of the label

To understand this, let assume  $n = 6$  (batch size) .

$y_1 = [1, 0, 0, 0], y_2 = [0, 1, 0, 0], y_3 = [0, 0, 1, 0], y_4 = [0, 0, 0, 1], y_5 = [1, 0, 0, 0], y_6 = [1, 0, 0, 0]$

And the output of NN is:

$\hat{y} =$

0.8	0.1	0.1	0
0	0.5	0.2	0.3
0	0	0.9	0.1
0.9	0	0	0.1
0.8	0	0.1	0.1
0.7	0.1	0.1	0.1


As it can be seen, the cross-entropy loss is minimized only if the NN predicts all the classes correctly.

Cross-entropy loss is a maximum likelihood estimator as opposed to MSE loss that we saw before.

To simplify the derivation, assume binary classification. For example, is this an image of a cat or not. This is also called **logistic regression**.

The NN has now only one output. And the output goes through a sigmoid function:

$$\hat{y} = \sigma(w^T x + b)$$



$$\text{sigmoid}(x) = \frac{1}{1 + \exp(-x)}.$$

$$P(y|x) = \begin{cases} \sigma(w^T x + b) & \text{if } y = 1 \\ 1 - \sigma(w^T x + b) & \text{if } y = 0 \end{cases}$$

Can be written in compact form:

$$p(y|x) = (\sigma(w^T x + b))^y (1 - \sigma(w^T x + b))^{1-y}$$

We want to maximize the probability of predicting the correct class ( $y$ ), given all  $n$  training inputs ( $x$ ).  
In other words, a classifier is a good classifier if it predicts **all the training examples** with high accuracy.

$$\begin{aligned} L(\mathbf{w}) &= P(\mathbf{y} \mid \mathbf{x}; \mathbf{w}) \\ &= \prod_{i=1}^n P(y^{(i)} \mid x^{(i)}; \mathbf{w}) \\ &= \prod_{i=1}^n \left( \sigma(z^{(i)}) \right)^{y^{(i)}} \left( 1 - \sigma(z^{(i)}) \right)^{1-y^{(i)}} \end{aligned}$$

$$\begin{aligned} l(\mathbf{w}) &= \log L(\mathbf{w}) \\ &= \sum_{i=1}^n \left[ y^{(i)} \log \left( \sigma(z^{(i)}) \right) + (1 - y^{(i)}) \log \left( 1 - \sigma(z^{(i)}) \right) \right] \end{aligned}$$



### Some final notes:

Pytorch has the cross-entropy loss and it automatically does the one-hot encoding of the class labels.

Logits:

refer to the raw, unnormalized scores output by a model before applying softmax.

We can show analytically that when taking derivative of the cross-entropy loss w.r.t the weights of the NN we get a clean expression for gradient:

$$\nabla_w L = -X^T(Y - A)^T, \quad W \in R^{k \times m}, X \in R^{n \times m}, A \in R^{n \times h}, Y \in R^{n \times h}$$

One-hidden layer.

*n: batch size*

*h: number of neurons in the hidden unit hidden*

*Y: onehot encoded matrix*

*m: input size*

But if we had used MSE loss, the gradient would have been very flat → difficult training.